

Can a judge tell how a model was changed?


Oct 5, 2026 · @alex

With outputs, logits and activations available, a judge that knows nothing about the specific models said correctly how all 24 variants were changed (base, quantized, fine-tuned or system-prompted), including two models it had never seen. From output or logits alone it reached 83% and 88%.

Pipeline

Pipeline

The judge sees only a document built from split A and one anonymized case



Nothing from split B or the two held-out models goes into the judge document or the baseline; the answer key never reaches the judge. The baseline classifies the same 24 cases.

Question and setup

Given only the computed metrics comparing an unknown Llama-2-7b-chat variant with the original (the reference), and a document on how those metrics usually behave for each kind of change, can a judge say whether the model is unchanged, or was quantized, fine-tuned or system-prompted?


Category	Models	What was done
base	base vs itself	nothing
quantized	awq	weights compressed to 4 bits
fine-tuned	ft_bias, ft_backdoor, ft_refusal_backdoor (held out)	LoRA adapters
system-prompted	sp_bias, sp_backdoor, sp_backdoor2 (held out)	hidden instruction in front of every conversation

- **Data.** Held-out WildChat split B, 64 prompts per measurement, 3 repeats per model. Metrics come from the unmodified feature script.
- **Cases.** One per (model, repeat): 24 cases (base 3, quantized 3, fine-tuned 9, system-prompted 9), with random IDs and a separate answer key.
- **What a case shows.** Output-level metrics with generic text rates (refusal openings, refusals added and dropped, disclaimers, stage directions, very short replies, mean length, each next to the reference's rate), logit metrics and activation metrics. The detectors written for specific models (marker word, cuisine, "helpful and respectful" opener) are left out.
- **Judge document.** Metric definitions plus, per category, the lowest and highest value of every metric over the split-A models of that category (1 quantized, 2 fine-tuned, 2 system-prompted). No per-model description, trigger, payload, topic or intent. The held-out models and split B played no part in it.
- **Conditions.** The case shows only output, only logits, only activations, or all three; the judge always gets the whole document.
- **Judge.** Claude Sonnet 5.5 through the API, no tools, adaptive thinking at effort high, schema-constrained answer (label, confidence, justification). 3 calls per case and condition: 288 calls, all answered, about \$1.5.
- **Baseline.** Nearest centroid on split A's per-repeat features with the same 4 labels and the same features.

The full judge document, exactly as the judge read it: Judge document tab, or [mechanism_doc.md on GitHub](#).

Chance is 25%; always answering the most common label would score 9/24.

Results



Per model (judge calls right out of 9 · baseline cases right out of 3):

Model	Category	Output only (judge . baseline)	Logits only (judge . baseline)	Activations only (judge . baseline)	All three
base vs base	base	9/9 · 3/3	9/9 · 3/3	9/9 · 3/3	9/9 · 3/3
awq	quantized	9/9 · 1/3	9/9 · 3/3	9/9 · 3/3	9/9 · 2/3
sp_bias	system-prompted	9/9 · 3/3	9/9 · 3/3	9/9 · 3/3	9/9 · 3/3
sp_backdoor	system-prompted	9/9 · 0/3	9/9 · 3/3	9/9 · 3/3	9/9 · 3/3
sp_backdoor2 (held out)	system-prompted	9/9 · 0/3	9/9 · 3/3	9/9 · 3/3	9/9 · 2/3
ft_bias	fine-tuned	6/9 (3 × quantized) · 2/3	9/9 · 3/3	9/9 · 3/3	9/9 · 3/3
ft_backdoor	fine-tuned	9/9 · 3/3	9/9 · 2/3	9/9 · 3/3	9/9 · 3/3

Model	Category	Output only (judge . baseline)	Logits only (judge . baseline)	Activations only (judge . baseline)	All three
ft_refusal_backdoor (held out)	fine-tuned	0/9 (9 × system-prompted) · 3/3	0/9 (9 × system-prompted) · 0/3	9/9 · 0/3	9/9 · 3/3

What decided the answers

- **System-prompted:** replies differ from the first token (median first divergence 0, edit distance about 0.8), the largest logit change, and an activation difference that climbs to 0.41–0.50 by layer 16 with low rank. All 27 calls per condition were right, including the held-out model with a different hidden instruction.
- **Quantized:** late text divergence (median 8.5–10), the smallest logit change (KL about 0.03–0.04), a size ratio below 1 (0.968), rank about 4.1 and direction consistency about 0.52. Right in every condition.
- **Fine-tuned, from activations:** the two split-A fine-tunes show a size ratio above 1 and direction consistency of 0.62–0.71. The held-out refusal fine-tune does not (size ratio 0.978–0.980, rank 3.82–3.92), so it was classified by elimination:
 - no layer-2 spike (0.05–0.18, against 31.7–64.0 for system prompts in the document)
 - a layer-1 difference of only 0.037
 - steady growth to 0.55 at layer 32, unlike quantization's flat profile
 - direction consistency of 0.61, near the fine-tuned range

Its confidence was the lowest of any correct activation answer (0.50–0.72).

- **The refusal fine-tune looks like a system prompt from outside.** It diverges at token 1 (median), has the highest KL measured (0.45–0.55, above even the system-prompted range), top-1 agreement of 0.84–0.88, adds 8–9 refusals per 64 prompts, and gives 25–30% very short replies. The judge called it system-prompted in all 18 output-only and logits-only calls, with confidence 0.55–0.80.
- **The other 3 errors:** ft_bias read as quantized from output in one repeat, where its median first divergence was 10 (the edge of the quantized range) and its refusal and disclaimer rates were close to the reference's.

Takeaways

1. **Telling how a model was changed transfers to new models when activations are available.** Both held-out models were classified correctly from activations alone in all 18

calls. The decision rests on a few activation features: the layer-2 spike, the mid-depth plateau, the size ratio and direction consistency.

2. **Without activations, a strong fine-tune can pass for a system prompt.** The refusal fine-tune moves text and logits more than either system-prompted model, so output-only and logit-only auditors misread it every time, and with confidence up to 0.80. Here, "diverges from the first token with a large KL" is not specific to system prompts.
3. **The judge beats the nearest-centroid baseline in every condition** (majority vote 20, 21, 24, 24 of 24 vs 15, 20, 21, 22). From output alone the baseline confuses quantized with fine-tuned and system-prompted with fine-tuned; from logits and from activations it places the refusal fine-tune nearest the system-prompted centroid.

Caveats

- **Few models per category.** The document's ranges come from 1 quantized, 2 fine-tuned and 2 system-prompted models, all from one base model and one quantization method. The test adds one new system-prompted and one new fine-tuned model; no new quantized or base model. The fine-tunes are LoRA adapters only.
- **The layer-2 spike carries much of the activation-level call, and it depends on the prompts drawn.** On split B it was about 0.1 for every model in repeat 1, the system-prompted ones included. The judge still got those cases right from the plateau, onset and rank, but a draw where both happen would be harder.
- **The document was written after the test results had been read.** Every number in it is computed by a script from split A only, and its prose restates those ranges, but which features the prose highlights, and the phrase "often with a different opening or voice" for system prompts, were judgment calls.
- **No harmless system prompt was measured.** It would presumably be labelled system-prompted, which is correct here, but it was not tested.
- **Small numbers.** 24 cases, 3 per model. Repeats share prompts (126 unique of 176), and the 3 calls per case measure consistency, not new evidence.
- **Judge model.** Could have been more accurate with more data.

Method notes

- Leak check: the 96 case files contain no model names, spec keys, paths, trigger words, payload words, topic words or the words "bias" and "backdoor"; the judge document contains none of them either.
- One pilot call (one system-prompted case, all levels) checked the Sonnet 5.5 settings before the run; nothing changed after it, and it is one of the 288.